



Anchors for Nuke — User Guide

Narrative walkthrough and complete reference · The anchors plugin

Contents

1	Anchors — User Guide	1
2	Anchors in practice	2
2.1	The problem: reusing a node	2
2.2	The solution: anchors and links	6
2.3	A real comp	8
2.4	Organising sources: the “input block”	10
2.5	Reusing and expanding	11
2.6	Finding your way	13
2.7	Anchors: the three tiers	16
2.7.1	Anchors	16
2.7.2	Dot anchors	17
2.7.3	Local Dots	18
2.8	Creating links	19
2.9	Copy, cut, and paste with hidden-input reconnection	19
2.10	Navigation	21
2.11	The spatial view	23
2.12	The leader key	24
2.13	Reconnecting links	24
2.14	Colours	25
2.15	Automatic naming	26
2.16	Label utilities	26
2.17	Setting up a backdrop	26
2.18	Upgrading a script built with another tool	27
2.19	Preferences and site configuration	28
2.19.1	Space-prefix search modes	28
2.20	Python API	28
2.21	Keyboard reference	29

1 Anchors — User Guide

Anchors is an alternative to systems like the popular *Stamps* library, creating a named anchor-and-link system for reusable inputs. Drop a named *anchor* on any stream, then reference it from lightweight *Link* nodes anywhere in the script with a colour-coded fuzzy find menu (like the Tab menu - accessed through “A”). Copy and paste become reconnection-aware, so pasted Links retain their inputs instead of arriving with dangling pipes.

Anchors has a couple main advantages over competing solutions. Firstly is performance. Links are native hidden inputs - no callbacks, no expressions, and therefore no performance impact on your script. Second is the built-in navigation workflow - you can jump to any Anchor

(or labelled backdrop) with a convenient fuzzy-find menu (Alt-A), and jump back to where you were (Alt-Z), which means as you build your script, it remains easy to navigate.

2 Anchors in practice

2.1 The problem: reusing a node

Every comp reuses the same source in different places, sometimes in many many places. A plate feeds the base of the comp and is reused many times in the foreground layer; a CG render feeds the beauty AND is needed in different places for Cryptomattes, P_Mattes, IDs, etc., and in the foreground section to limit the range of edge work; a camera drives multiple different projections and renders.

Nuke gives you two obvious ways to reuse a node, and both go wrong as the script grows.

Duplicate the node. Copy the Read and paste it wherever you need it again. Now the same logical source lives in the script multiple times (Figure 1), and every one of those copies is a separate node to keep in sync. With the best will in the world, in a complex script, mistakes will be made when it comes time to update the node. It's all too easy to update 11 of the 12 places the CG layer is used and introduce a subtle bug into your comp.




Figure 1: Reuse by duplication: the same plate and CG Reads copied all over the script — every copy is another node to keep in sync.

Run a long pipe. Keep a single Read and drag its output across the whole script to each place that needs it (Figure 2). Now there is only one node to maintain, but the graph is a cat's cradle of long pipes: these are extremely hard to read, very easy to disconnect by accident, and awkward to lay out — every new node means routing around the pipes that already cross the script, which makes working in the script frustrating and error prone, and makes one shy away from the kinds of restructuring that scripts often need.




Figure 2: Reuse by long pipes: one Read, but its output is dragged the length of the script. Hard to maintain, miserable to read.

A third option — a bare hidden input — keeps the graph clean until you copy/paste a block that contains it, at which point the hidden connection silently breaks and leaves a node wired to nothing, with no record of what it was meant to connect to.

2.2 The solution: anchors and links

An *anchor* is a small named node you drop onto a stream. A *link* is a lightweight node that references an anchor BY NAME — its input is hidden, so the connection travels by name instead of a drawn pipe. Wherever you would have duplicated a Read or run a long pipe, you drop one anchor on the source and place links instead (Figure 3).




Figure 3: The same comp with anchors: one `Anchor_plate` and one `Anchor_cg` at the top, and lightweight `Link:` nodes wherever the plate or CG is reused. One node to maintain, and no long pipes.

You get the tidy graph of a hidden input AND the safety of a named reference: the plate and the CG each exist exactly once, the links carry their names so the graph reads at a glance, and — as we will see — a block of links repairs its own connections when you copy and paste it, even into another script. The script is easy to understand, easy to work on, and easy to expand. Blissful productivity.

2.3 A real comp

Here is the same idea demonstrated in a simple comp script (Figure 4): a CG render with breakout and mattes, a projected matte painting, a key, and a rotoed core matte. The sources are collected into an *input block* at the top of the script, and the comp below is a single clean spine of links — no source is ever piped across the graph or duplicated.



Figure 4: A small example comp: the input block up top (all sources: plates, CG, matte-painting, cameras, geo), and readable modules below, each with its own clear B-spine.

2.4 Organising sources: the “input block”

This is the idiomatic usage of Anchors, and is the absolute best way to build a script. All sources live in one place at the top of the script with an Anchor attached to allow Links to them throughout the script (Figure 5).

Colour groups them by type — plates gold, CG and matte-paintings blue, cameras and geo green — so the whole cast of the comp is legible in one spot, and every link elsewhere in the script inherits its anchor’s colour.

Of course, the colours are up to you, and if - as in this example - you place each type in a coloured backdrop, the Anchors will by default pick up the colour of the backdrop, making it easy to be consistent.




Figure 5: An input block: each source anchored and colour-coded by type, forming a tidy library the rest of the comp links to.

2.5 Reusing and expanding

With the sources anchored, extending the comp is easy. Working from the example above, we can add a clean-plate projection setup to the IBKGizmo.

Hit “A” to bring up the fuzzy-find menu (Figure 6) and drop links to the camera and the plate and wire them up — no new Reads, no pipes dragged across the script (Figure 7). And because your module inputs have clean hidden inputs, your modules remain visually and logically distinct from the rest of the script and easy to enlarge with more nodes.




Figure 6: The Anchor selection menu.




Figure 7: Building a new branch by linking to existing anchors — the camera, plate, and clean-plate are all links, not new Reads.

The same anchors serve a second version of the comp just as easily — a variant key and a keymix built entirely from links to the sources you already anchored (Figure 8).




Figure 8: Reusing the same anchors for a second treatment: a fresh key and keymix, all driven by links.

2.6 Finding your way

One feature we haven't touched on is *Anchor Dots*. Dots at the output of a module are the best way of labelling a module (sticky notes can get lost as the script is expanded/reorganized because they're not wired in, backdrops create busy work expanding the backdrop as you add nodes).

Anchors takes advantage of and enables that workflow in two ways. One, Shift-B/N/M respectively label a Dot with a small/medium/large label (perfect for labelling modules/sub-modules/sub-sub-modules and making clear that hierarchy). Two, Dots labelled in this way are *Anchor Dots* - they get a special colour and operate the same as regular Anchors in your script.

As a script fills with links, the anchors become navigation landmarks. Press **Alt+A** (Figure 9) for a fuzzy-search picker of every anchor and labelled backdrop and the graph zooms straight to the one you choose (Figure 10). With a link selected, **Alt+J** jumps to its source anchor (Figure 11).

Either way the graph frames the anchor together with the tree feeding it, so you land on

the whole module. When you would rather an anchor were ONLY a landmark — a bookmark part-way down a long spine, say — tick *Jump to anchor only* in its properties panel and jumps to it centre on the anchor itself and leave its tree alone. The box is unticked on every existing anchor, so nothing changes until you ask for it.

After jumping with *Alt+A* or *Alt+J*, *Alt+Z* jumps you back to where you were.

This enables a wonderfully quick way of getting around your script. As you document your work with labelled dots, you're also creating a map of the script for quick navigation. Turn on the *spatial view* and the menu shows you that map directly — see [The spatial view](#).




Figure 9: The navigation picker (*Alt+A*).




Figure 10: Navigating to a module: the graph zooms to the matte-painting projection behind its dmp bg anchor.




Figure 11: Jumping from a link to its source: the `cg_env` anchor under its CG Read.

2.7 Anchors: the three tiers

2.7.1 Anchors

The standard anchor. Select a node and press *A* (or *Edit > Anchors > Create Anchor*). A NoOp anchor is created beneath the node, wired to it, and named from the source (Figure 12).

A matching *link* is created directly below the new anchor, so you get both halves of the pair from one gesture — drag the link wherever you need it. If you would rather create links yourself, switch off *Create a link below each new anchor* in Preferences.

Each NoOp anchor’s properties panel carries three buttons and a checkbox:

- *Reconnect Child Links* — rewire every link that points at this anchor.
- *Rename* — rename the anchor and update all of its links automatically.
- *Set Color* — open the colour palette and recolour the anchor and its links.
- *Jump to anchor only* — when ticked, jumping to this anchor frames the anchor by itself instead of the tree above it.

By convention, “true” Anchors are used for the inputs of the script, not for sections within the script.




Figure 12: An anchor pointing to a source node.

2.7.2 Dot anchors

For lightweight, in-graph signposting, promote a plain *Dot* to an anchor by selecting it and pressing a label key:

Key	Size	Font
Shift+B	Small	33 pt
Shift+N	Medium	66 pt
Shift+M	Large	111 pt

Any Dot whose label is *33 pt or larger* is treated as an anchor and appears in the navigation picker. Dot anchors are always the default purple and propagate their label to any links pointing at them (Figure 13). They carry the same *Jump to anchor only* checkbox as NoOp anchors.

By convention, *Dot Anchors* are used WITHIN THE SCRIPT, never for input sources.




Figure 13: Dot Anchors of various sizes used to identify and bookmark parts of a script.

2.7.3 Local Dots

A *Local Dot* is a hidden-input Dot that points at a PLAIN node (not an anchor). It reconnects only within the same script — never across scripts — and is stamped burnt-orange with a **Local:** <source> label. Local Dots are produced when you hide the input of a wired Dot (*Alt+H*, or *Edit > Node > Input On/Off*).

These are handy for re-using sections of a module within the module. If you're re-using something across modules, it's better to create a *Dot Anchor* as a named source.

2.8 Creating links

Unless you have switched the preference off, every new anchor already arrives with a link beneath it (see *Anchors* above). To create further links to an existing anchor:

With *nothing selected*, press **A** to open the Anchor selection menu. It lists every anchor (Figure 14) in the script with its colour; choose one and a link is created at the cursor, wired to that anchor.




Figure 14: The Anchor selection menu.

The link type is chosen automatically from the source: a *Dot* source produces a *Dot* link; anything else produces a *NoOp* link. Each link carries a *Reconnect* button and inherits the anchor's colour, updating automatically when the anchor's colour changes.

You can also create links for several selected anchors at once via *Edit > Anchors > Create Link*.

2.9 Copy, cut, and paste with hidden-input reconnection

The plugin overrides the standard clipboard shortcuts (scoped to the Node Graph):

- **Ctrl+C** — copy. Hidden reference knobs are stamped on the selected nodes before copying, then the originals are restored (the copy is non-destructive).
- **Ctrl+X** — cut. As copy, but moves anchors cleanly.
- **Ctrl+V** — paste. Each Link looks up its anchor (Figures 15 and 16) and rewires:
 - *Same script, same scope*: links reconnect to the original anchors.

- *Across scripts*: links reconnect by the anchor's display name if a matching anchor exists in the destination.
- *Local Dots*: reconnect by identity, same-script only.
- *Pasting only anchors*: If you copy ONLY Anchors, each pasted anchor is replaced by a Link pointing back to the original, so you do not get duplicate anchors.




Figure 15: A reusable block containing two Links, before copying.




Figure 16: After pasting, both Links have reconnected themselves to their anchors.

Paste Multiple (*Edit > Paste Multiple*) pastes repeatedly, re-piping each copy. For a plain paste with no reconnection magic, use *Paste (old)* (*Ctrl+Shift+D*); *Copy (old)* and *Cut (old)* are also available under *Edit > Anchors*.

2.10 Navigation

Shortcut	Action
Alt+A	<i>Anchor Find</i> — fuzzy-search picker (Figure 17) of all anchors and labelled backdrops; zooms to the chosen entry.
Alt+J	<i>Anchor Jump</i> — with a Link selected, jump to its source anchor.
Alt+L	<i>Cycle Links</i> — with an anchor selected, step through each of its links. Keep pressing L to advance; any other key stops.
Alt+Z	<i>Anchor Back</i> — return to the viewport position from before the last jump.

Both *Anchor Find* and *Anchor Jump* frame the anchor and the tree feeding it. Tick *Jump to anchor only* on an anchor to make jumps to it land on the anchor alone — useful for anchors that are navigation landmarks rather than the head of a module. The setting lives on the anchor, so it applies however you jump to it.




Figure 17: The navigation picker (Alt+A).

2.11 The spatial view

Sometimes you know WHERE an anchor is without remembering what it is called. Tick *Show the spatial view beside the A and Alt+A menus* in Preferences and both menus — and the leader's *Set Input* commands, which use the same anchor menu — open with a map of the script beside their search list (Figure 18). The search field and list work exactly as before; the map is an addition.




Figure 18: The spatial view beside the Alt+A menu.

The map draws each item where it really sits in the DAG, with the empty space between items squeezed out, so what is left of what and what is above what stays the way you remember it:

- *Anchors* are coloured tiles.
- *Dot anchors* are circles with their label beside them, in three sizes that follow the Dot's label size — small, medium, or large, whichever of the **Shift+B / N / M** presets it is nearest.
- *Labelled backdrops* are frames drawn around the anchors, Dot anchors, and other backdrops they enclose in the DAG. A backdrop with no anchors inside it is drawn as a faded, dashed box so it never reads as an anchor.
- *Local Dots* are not shown.

Items are packed only as close as they can come without overlapping: the empty space shrinks, but nothing changes places — a module to the right of another in the DAG is to its right on the map, and one above another stays above it.

As you type, the list filters as usual and the map items that no longer match grey out rather than disappearing, so the map keeps its shape while you narrow it down. The entry highlighted in the list is outlined in white on the map, so **Up / Down** walk the map too. Click any map item to pick it, exactly as if you had chosen its row. In the **A** menu the backdrops are drawn for context only: you link to anchors, so only the anchors can be clicked.

In a big script the map does not fit beside the list at a readable size, so the map zooms to what you are looking for instead. As you type, it zooms and scrolls to frame the items that still match, and it never shrinks the map past the point where labels stop being legible. You can take over whenever you want, with the same mouse moves as the Node Graph: the mouse wheel zooms in and out around the pointer, and dragging with the middle button moves the map. The zoom controls in the map's bottom-right corner zoom too: **-** and **+** (or **Ctrl+- / Ctrl+=**) zoom out and in. Once you have zoomed, the zoom stays put and typing only scrolls the map when the highlighted entry is out of sight, bringing it to the middle. *Fit* (or **Ctrl+0**) hands the zoom back to the search, and so does opening the menu again. Clicking the controls

leaves the search field focused, so you can keep typing. Every zoom and scroll glides into place rather than jumping, so you can follow where it went. While the map is bigger than its area, a *minimap* above the zoom controls shows the whole script with the visible part outlined; click or drag in it to move there.

With the preference off (the default), these menus are the plain lists shown elsewhere in this guide.

2.12 The leader key

Press **Shift+A** to raise the *leader overlay* — a translucent heads-up map of every anchor command (Figure 19). Press one key to run the action; the overlay then disappears. Cells grey out when their preconditions are not met (for example, **W** is disabled on a single-input node).




Figure 19: The leader overlay maps every command to a single key.

Key	Action
Q	Set <i>B</i> input from an anchor (input 0)
W	Set <i>A</i> input from an anchor (input 1)
E	Set <i>Mask</i> input from an anchor
R	Set the <i>first free</i> input from an anchor
F	Anchor Find (as Alt+A)
J	Anchor Jump (as Alt+J)
L	Cycle Links (as Alt+L ; keep pressing to chain)
Z	Anchor Back (as Alt+Z)
X	Reconnect All Links
,	Open Preferences

The overlay grid matches your physical keyboard; *QWERTY*, *AZERTY*, and *QWERTZ* layouts are supported (set in Preferences).

2.13 Reconnecting links

Links are stateless — they resolve their anchor on demand — so you can rewire at any time:

- *Reconnect All Links* (*Edit > Anchors > Reconnect All Links*, or leader **X**) rewires every link in the script. Handy after loading or merging scripts.
- *Reconnect Child Links* (button on each anchor) rewires just that anchor's links.
- *Reconnect* (button on each link) rewires that single link.

Because there are no callbacks or other shenanigans, nothing is stopping you from disconnecting a Link node, but you can always reconnect it if it does get disconnected.

2.14 Colours

Anchors carry a tile colour (Figure 20) so related streams are easy to spot at a glance.




Figure 20: Anchors carry colours so related streams stand out.

Open the colour palette (Figure 21) from an anchor's *Set Color* button. It offers Nuke's preference colours, the colours already used by backdrops in the script, and your own saved palette; *Custom Color...* opens a full picker.




Figure 21: The colour palette.

Picking a colour applies it and closes the palette straight away. If you would rather compare a few colours before committing, untick *Selecting a color closes the color palette* in Preferences: the palette then only highlights each colour you pick and stays open until you confirm with **Enter** or **OK** (or discard it with **Esc**).

When you create an anchor, the dialog also offers a name field (Figure 22) so you can name and colour it in one step. New anchors pick a colour that contrasts with their containing backdrop, so an anchor inside a coloured backdrop stays legible. Dot anchors are always the default purple. Changing an anchor's colour updates all of its links.




Figure 22: Creating an anchor: name and colour together.

2.15 Automatic naming

When you create an anchor, the plugin suggests a name derived from the source — typically from the read file path and the containing backdrop. The naming rule is configurable in Preferences (an advanced *regex* and *template*, with a demo filename to preview the result), so a studio can standardise anchor names. Special characters are sanitised automatically; a name that sanitises to nothing is rejected.

2.16 Label utilities

Beyond Dot promotion, the label keys are general-purpose (they act on the selected node):

Key	Action
Shift+M	Label (Large) — 111 pt on Dots (promotes a plain Dot to an anchor), 33 pt on other nodes
Shift+N	Label (Medium) — 66 pt on Dots
Shift+B	Label (Small) — 33 pt on Dots

2.17 Setting up a backdrop

Backdrops carry the comp’s structure, so **A** treats them as something to set up rather than something to anchor. Select a single backdrop, press **A** (or use *Edit > Anchors > Setup Backdrop*), and a dialog opens with everything a backdrop needs in one place (Figure 23):

- a *label* field — multi-line, so a backdrop’s notes survive a round trip;
- the same *colour palette* as the anchor dialogs, including your custom colours and hint-mode navigation;
- a *font size* dropdown offering the Dot-anchor sizes — Small (33), Medium (66), Large (111) — plus *Custom* for anything else. A backdrop already set to one of the three presets re-opens on that size; anything else, including Nuke’s own default, opens on Large;
- a *Filled* checkbox — off draws the backdrop as an outline only.

Clicking a swatch confirms the whole dialog, as it does when renaming an anchor, so a label and a colour are two keystrokes and a click. Pressing **Enter** inside the multi-line label field adds a line; **Ctrl+Enter**, the *OK* button, or a swatch click confirms. Selecting a backdrop TOGETHER WITH its contents keeps the old behaviour and creates an anchor from the selection.




Figure 23: Setting up a backdrop: label, colour, font size, and fill in one dialog.

2.18 Upgrading a script built with another tool

Plenty of scripts already contain an anchor rig built by a different tool: a labelled, coloured NoOp under a Read, with hidden-input nodes dotted around the comp pointing back at it. It is the same idea as anchors, without the anchor machinery — so the picker, navigation, reconnect and colour propagation all ignore it.

Edit > Anchors > Upgrade to Anchors... adopts that rig. Each parent node becomes a real anchor, and every hidden-input node pointing at it becomes a real Link. Nodes are converted *in place*: a PostageStamp stays a PostageStamp, and every node keeps its position and its downstream connections.

The dialog previews exactly what will change before anything is touched, and offers:

- *Scope* — the selected nodes, or every anchor-like node in the script.
- *Parent nodes to upgrade* — NoOp and PostageStamp parents, Dot parents, or both. They are listed separately because the two usually want different naming.
- *Anchor names* — take the name from the node's label, its node name, or (the default) its label falling back to its node name. Set separately for NoOp and Dot parents, since a foreign NoOp usually carries a meaningful node name while a Dot is called something like `Dot17` and keeps its meaning in the label.
- *Strip leading / trailing text* — drop a tool's fixed affix, turning `Pointer_Foo` into `Foo`. A strip that would leave nothing behind is ignored.
- *Colours* — keep each node's existing tile colour, or take the colour the plugin would derive for a new anchor. Dot anchors always take the default anchor colour, as they do everywhere else.

Names are sanitised and made unique, so two parents that reduce to the same name become `Anchor_Foo` and `Anchor_Foo1`. A node that is BOTH a parent and someone else's hidden-input child stays a parent — it becomes an anchor rather than a Link. Nodes that are already anchors keep their name and colour; only their children are upgraded. Running the upgrade a second time does nothing.

Like the other migrators, this is not undoable — save a backup of your script first.

2.19 Preferences and site configuration

Edit > Anchors > Anchor Preferences... controls:

- *Enable anchors plugin* — the master toggle.
- *Create a link below each new anchor* — on by default; uncheck it if you prefer creating links yourself.
- *Show the spatial view beside the A and Alt+A menus* — off by default; tick it to add a map of the script to the anchor menus (see [The spatial view](#)).
- *Keyboard layout* — QWERTY / AZERTY / QWERTZ for the leader overlay.
- *Space-prefix search modes* — what leading spaces do in the fuzzy-find menus (see below).
- *Custom Colors* — add, edit, and remove the colours in your personal palette.
- *Selecting a color closes the color palette* — on by default; turn it off to keep the palette open until you confirm with **Enter** or **OK**.
- *Advanced* — the anchor *naming regex* and *template*, plus a *site config override*.

Settings are stored per user in `~/.nuke/anchors_prefs.json`. A studio can *publish* the naming settings to a shared site-config file (pointed to by the `ANCHORS_SITE_CONFIG` environment variable) so every artist gets consistent anchor names; individuals can opt out with the site-config override.

2.19.1 Space-prefix search modes

The **A** and **Alt+A** menus filter as you type, and typing one or two spaces before your search text switches how that text is matched:

- *Anchored fuzzy* — the letters appear in order, starting at the first letter of the name: `bgp` finds `BG_Plate`.
- *Non-anchored fuzzy* — the letters appear in order anywhere in the name: `plt` finds `BG_Plate`.
- *Consecutive substring* — the text appears as typed: `plat` finds `BG_Plate`.

By default no leading space means anchored fuzzy, one space means non-anchored fuzzy, and two spaces mean a consecutive substring search. The *Space-prefix search modes* group in Preferences maps each of the three levels to whichever mode you prefer; every mode must be used exactly once, so OK is refused if you assign the same mode twice.

These menus share their search core with [tabtabtab-nuke](#), which offers the same preference. If you run both, tick *Use tabtabtab-nuke preferences* and anchors follows the mapping set there instead — the three dropdowns grey out and show what tabtabtab is enforcing, and your own mapping is kept for when you untick the box. The checkbox itself is greyed out when no tabtabtab-nuke installation is found.

2.20 Python API

For pipeline and templating tools, `api.py` is the stable, documented surface. Both functions require a running Nuke session.

```
from api import create_anchor, find_anchor_by_name

# CREATE AN ANCHOR WIRED TO A SOURCE NODE, WITH AN EXPLICIT COLOUR.
source = nuke.toNode('Read_BG')
anchor_node = create_anchor('BG_Plate', input_node=source, color=0x8040FFFF)

# LOOK ONE UP BEFORE CREATING A DUPLICATE.
existing = find_anchor_by_name('BG_Plate')
if existing is None:
    existing = create_anchor('BG_Plate')
```

- `create_anchor(name, input_node=None, color=None)` -> the new anchor node. Raises `ValueError` if NAME sanitises to empty.
- `find_anchor_by_name(display_name)` -> the matching anchor node, or `None`.

2.21 Keyboard reference

Shortcut	Action
A	Create anchor (selection) / Set up backdrop (one backdrop selected) / Anchor selection menu (no selection)
Shift+A	Leader-key overlay
Alt+A	Anchor Find / navigate
Alt+J	Anchor Jump (Link -> anchor)
Alt+L	Cycle Links
Alt+Z	Anchor Back
Alt+H	Toggle input visibility (make a Local Dot)
Ctrl+C / Ctrl+X / Ctrl+V	Copy / Cut / Paste with reconnection
Ctrl+Shift+D	Paste (old) — plain paste, no reconnection
Shift+M / Shift+N / Shift+B	Label Large / Medium / Small
Ctrl+M	Append to label

All anchor shortcuts are scoped to the Node Graph, so they never intercept keys in the Viewer or Script Editor.